

HW based Performance Analysis Methodology

D. Levinthal

Overview

- Objectives
- Cycle Accounting
- Measuring performance limitations
 - Load latency
 - Memory bandwidth
 - Instruction starvation
 - Branch misprediction
 - Branch execution (calls and indirect branches)
 - Store resource limitations
 - Profiling and active binary size
- Microkernels and documentation
- Existing tools
- Actually running an application and collecting data
 - And common pitfalls

Performance analysis objectives

- Most performance analysis infrastructure developments driven by desire to improve code performance
 - Ability to identify largest easily fixed issue
 - Reduce execution cycle counts
- Performance event rates/instruction retired can guide processor design in simulators running traces
- Understand utilization at scale to identify deficiencies and guide requirements for next generation systems
 - CPU, mobo, addin cards

What matters is cycles

- Performance events count occurrences of signals in the processor
- These need to be turned into a single figure of merit for relative importance evaluation
- Cycles are the figure of merit that users care about
- Thus performance events that can be turned into cycle counts are the most valuable
 - Cache/dram hits can be multiplied by latencies
 - Cache misses cannot and ultimately tell you little
 - Profiling differences just does not work at all
- Dominant loss of performance on servers is due to stalled cycles
 - Cycles with zero instructions (issued, executed or retired)

| | |
|------------------------------|--------|
| IPC | 0.7428 |
| instructions/unstalled_cycle | 2.776 |
| uops(fused)/unstalled_cycle | 3.000 |

What matters is cycles

- SUM of event_counts*penalties is essentially serializing “bottleneck model”
- OOO execution will result in overlapping issues
 - Penalties can be greater than total cycles
 - Need to see all issues to identify which to attack
 - Use multiple metrics for consistency checks when possible
- Forcing a sum to 1.0 is intrinsically misleading
 - And just a bad idea
- **Any hw event based micro architectural performance decomposition methodology should consider achieving a 10% accuracy in its components a triumph**
 - **Don't expect better**

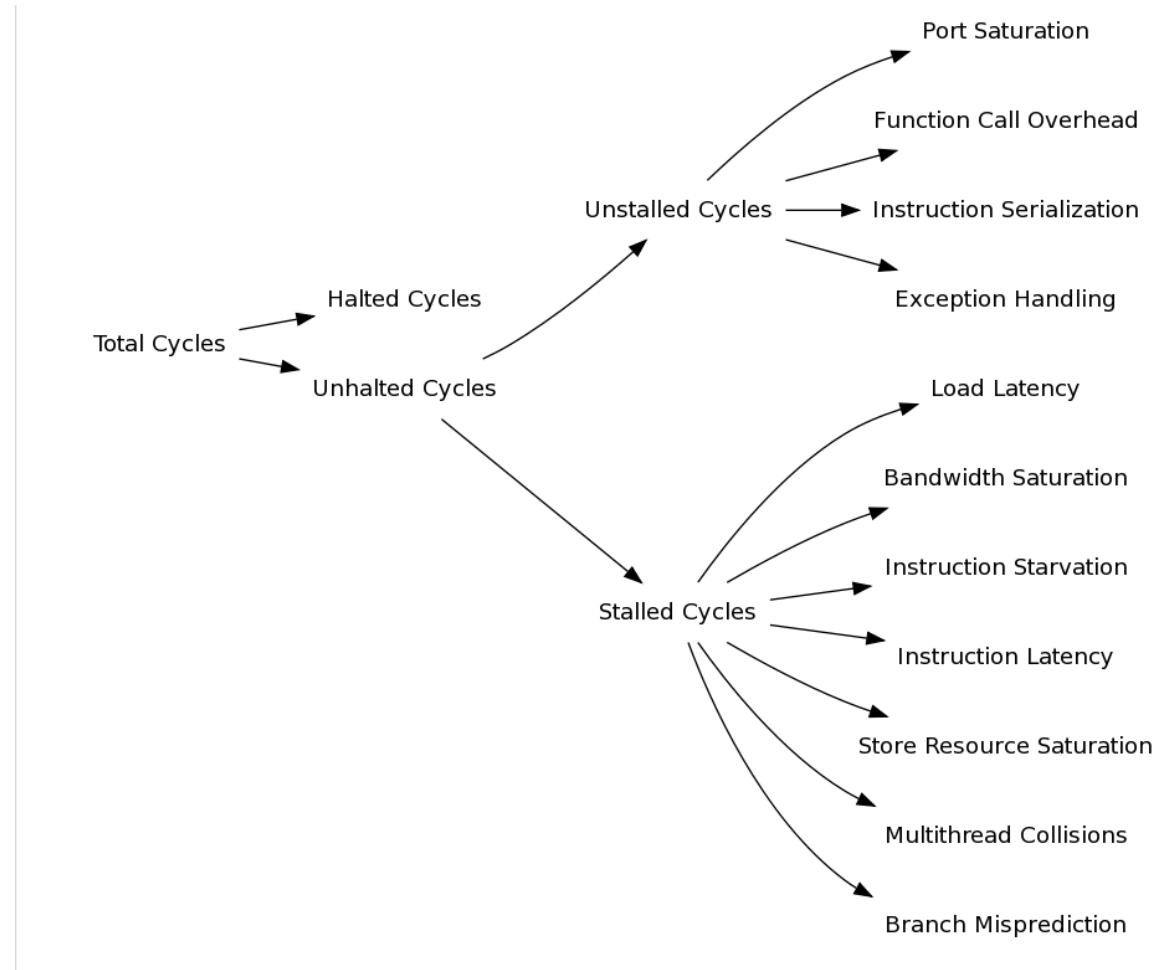
Identifying HW performance limitations in the Hardware

- Execution Characterization
 - Simple counting mode analysis based on well defined and validated performance events
 - Assign performance cost to architectural components
 - Latencies
 - resource saturation
 - Execution path discontinuities
- Can be done for a single application
- Or an entire data center
 - Statistical distributions

Identifying SW performance limitations in a single application

- Execution time is the figure of merit
- Improving that is made easier by precise knowledge of
 - Where in the code the time is spent
 - Why it is taking so long
- Standard approach is profiling code versus many performance events
- Event counters programmed to interrupt on overflow
 - ~ 1-4 khz interrupt rate/core results in a few % perf degradation
 - Interrupt data collected on disk
 - post process data with binary, source and debug symbol file

Cycle accounting and the Gooda cycle tree



**The leafs are constructed from bottom up sums, all using consistent units (cycles)
Consequently, drilling down to lower dominant components will always show
consistent relative values**

Generic Cycle accounting

| pipeline status | issue | estimating event description |
|-----------------|---------------------------|--|
| stall | load latency | number of cycles with pipeline stalls at execution stage while memory subsystem has an outstanding load |
| stall | bandwidth saturation | number of cycles where # of outstanding memory requests per core (not per thread) is near the minimum number measured at steady state with any number of triads running |
| stall | instruction starvation | number of cycles waiting at the execution stage for long latency code delivery (> 40 cycles). Also requiring the Scheduler/ROB being nearly empty |
| stall | instruction latency | Cycle count for divides, SQRT, transcendentals and other long latency execution streams not covered otherwise. |
| stall | store resource saturation | number of cycles all store buffers in use |
| stall | branch misprediction | total cost in cycles of branch mispredictions, including all time spent fetching/executing the wrong path, plus extra cycles needed to recover and resume progress on the correct path |
| stall | multithread collision | |
| unstalled | port saturation | cycles where single execution port getting uops nearly every cycle |
| unstalled | function call overhead | enable an estimated cost in cycles for calls including trampolines, stack manipulations, etc. This is done by using the metric instructions_retired/call_retired |
| unstalled | instruction serialization | cycles only executing chained dependent multi cycle latency uops |
| unstalled | exception handling | cycles spent issuing uops from a microcode scheduler for exception handling |

Load latency


```
<metric name="HIGH_LEVEL_CYCLE load latency penalty">
  <event alias="a">CPU_CLK_UNHALTED.REF_TSC</event>
  <event alias="b">MEM_LOAD_L3_MISS_RETIRED.LOCAL_DRAM</event>
  <event alias="c">MEM_LOAD_L3_MISS_RETIRED.REMOTE_DRAM</event>
  <event alias="d">MEM_LOAD_L3_MISS_RETIRED.REMOTE_HITM</event>
  <event alias="e">MEM_LOAD_L3_MISS_RETIRED.REMOTE_FWD</event>
  <event alias="f">MEM_LOAD_L3_HIT_RETIRED.XSNP_HITM</event>
  <event alias="g">MEM_LOAD_L3_HIT_RETIRED.XSNP_HIT</event>
  <event alias="l">MEM_LOAD_L3_HIT_RETIRED.XSNP_MISS</event>
  <event alias="h">MEM_LOAD_RETIRED.L3_HIT</event>
  <event alias="i">DTLB_LOAD_MISSES.WALK_COMPLETED</event>
  <event alias="j">DTLB_LOAD_MISSES.WALK_ACTIVE</event>
  <event alias="k">CPU_CLK_UNHALTED.THREAD</event>
  <formula>((167*b+368*c+335*(d+e))/a+(189*(f+g+l)+75*h+3.5*i+j)/k)</formula>
</metric>
```

Memory Bandwidth

- Memory bandwidth is a problem only when you get close to the limit
- Bandwidth is a time average
 - If the time scale of the saturation is too small you cannot detect the saturation by measuring BW
 - Ex: 10% of every 1 msec is running memcpy
 - Measuring BW would indicate you are using only 10% of available BW
 - No problem
- One needs to measure on each cycle if the BW is saturated
- This is done by monitoring cacheline request queue occupancies
- Punchline: ≥ 6 offcore_requests in the offcore queue is a saturated cycle

Memory BW saturation on CLX using parallel triads BW

Total BW varies from 19GB/sec to 95



This approach can identify short durations BW saturations
But also even if only a single core is BW limited






Branch misprediction

```
<metric name="HIGH_LEVEL_CYCLE branch misprediction penalty">  
  <event alias="a">CPU_CLK_UNHALTED.THREAD</event>  
  <event alias="b">BACLEAR.S.ANY</event>  
  <event alias="c">BR_MISP_RETIRED.ALL_BRANCHES</event>  
  <event alias="d">INT_MISC.RECOVERY_CYCLES</event>  
  <event alias="e">UOPS_ISSUED.ANY</event>  
  <event alias="f">UOPS_RETIRED.RETIRE_SLOTS</event>  
  <event alias="g">UOPS_ISSUED.ANY:c1</event>  
  <formula>((6*(b+c)+d+((e-f)*g/e))/a)</formula>  
</metric>
```

Term $(e-f)*g/e$ is cycles spent issuing uops for mispredicted path

Instruction starvation/Frontend stalls

- OOO structures are huge on modern processors
 - FE not delivering uops in a few slots or even cycles just doesn't matter on servers
- If the scheduler has uops there is no problem
- So, only Ifetch from L3 is a real performance impact

| | Sandy Bridge | Haswell | SkyLake |
|-----------------------|--------------|---------|---|
| Out-of-order Window | 168 | 192 | 224  |
| In-flight Loads | 64 | 72 | 72 |
| In-flight Stores | 36 | 42 | 56  |
| Scheduler Entries | 54 | 60 | 97  |
| Integer Register File | 160 | 168 | 180  |
| FP Register File | 144 | 168 | 168 |
| Allocation Queue | 28/thread | 56 | 64/thread  |

Instruction Starvation/FE stalls

```
<metric name="HIGH_LEVEL_CYCLE instruction starvation cycle fraction">  
  <event alias="b">CPU_CLK_UNHALTED.THREAD</event>  
  <event alias="d">RS_EVENTS.EMPTY_CYCLES</event>  
  <formula>(d/b)</formula>  
</metric>
```

```
<metric name="metric_cycle L3_ifetch cycle fraction">  
  <event alias="b">CPU_CLK_UNHALTED.THREAD</event>  
  <event alias="d">FRONTEND_RETIRED.L2_MISS</event>  
  <formula>65*(d/b)</formula>  
</metric>
```

Store resource limitations

- There are a lot of store buffers (35-60) but servers still run out due to long latency of RFO from dram
 - Stores MUST finish in order
- FE blocks on first store after running out of store buffers

<metric name="HIGH_LEVEL_CYCLE no store buffer cycle fraction">

 <event alias="b">CPU_CLK_UNHALTED.THREAD</event>

 <event alias="d">RESOURCE_STALLS.SB</event>

 <formula>(d/b)</formula>

</metric>

Branch misprediction

- **total cost in cycles of branch mispredictions, including all time spent fetching/executing the wrong path, plus extra cycles needed to recover and resume progress on the correct path**

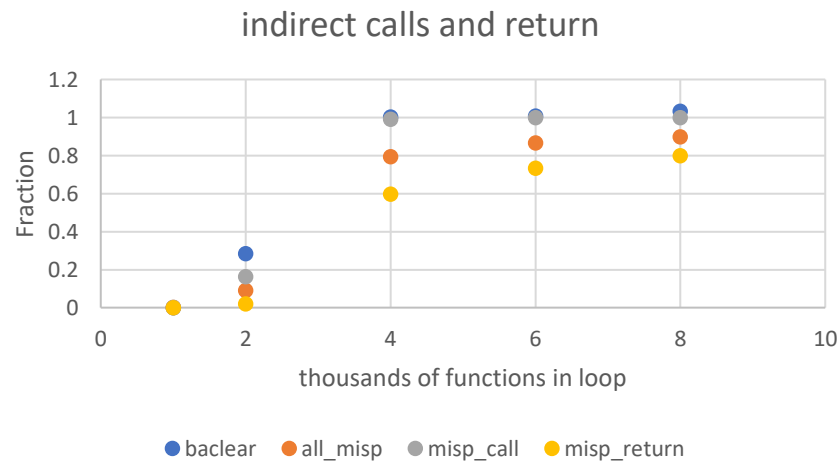


```
<metric name="HIGH_LEVEL_CYCLE branch misprediction penalty">  
  <event alias="a">CPU_CLK_UNHALTED.THREAD</event>  
  <event alias="b">BACLEARS.ANY</event>  
  <event alias="c">BR_MISP_RETIRE.ALL_BRANCHES</event>  
  <event alias="d">INT_MISC.RECOVERY_CYCLES</event>  
  <event alias="e">UOPS_ISSUED.ANY</event>  
  <event alias="f">UOPS_RETIRE.RETIRE_SLOTS</event>  
  <event alias="g">UOPS_ISSUED.ANY:c1</event>  
  <formula>((6*(b+c)+d+((e-f)*g/e))/a)</formula>  
</metric>
```

UOPS_ISSUED.ANY- UOPS_RETIRE.RETIRE_SLOTS is the mispredicted path length

UOPS_ISSUED.ANY:c1/ UOPS_ISSUED.ANY = cycles/issued uop

Branch Target Buffer and indirect branch misprediction

- Indirect branches can mispredict the target if the BTB is overwhelmed by the sheer number of branches in a large binary



Branch Target Buffer and indirect branch misprediction

- Indirect branches can cause a high fraction of branch mispredictions and L3 instruction fetching (<https://research.google/pubs/pub48320/>)
- The data just shown would suggest the use of a single common BTB.
- There might be benefit in having separate BTBs for branches that mispredict if there is no BTB target
 - I.e indirect branches and returns

Calling convention issues/ call overhead

- X86 calling conventions require arguments be put in particular registers
 - This can require instructions to push current values to the stack and replace them with call arguments
 - Calling a function in a shared object/DLL will result in trampoline branches
 - On arriving in the function, registers must be made available for the functions work. More registers pushed to the stack
 - Refactoring results in function argument validity being tested after entry
 - If this fails all the above work must be undone, a branch to recovery code and a return
- A call can easily result in an overhead of 10 -15 instructions
- Small functions can be costly
 - measured by `instructions_retired.any/br_inst_retired.any`

Deep call stacks


- Return instructions are indirect branches
- Deep call stacks can result in
 - mispredicted returns
 - Larger expense in restoring state
- Can be evaluated with `br_misp_retired.near_return`
 - If it is available

Profiling for HW evaluation

- Counter overflow driven interrupt sampling can be done with all the events discussed.
- However, server codes tend to have large active binary cores
 - A 100 MB binary can easily have 5 MB of active core
 - To sum to 75% of cycle interrupt samples likely requires 500-1000 functions
 - Hottest functions are likely malloc and free
 - Recall the google “data center tax”
- Sampling on instructions retired (or better LBRs) can yield basic block execution count
- Easier is simply counting instruction cacheline hit frequency
 - Yielding active binary size

Active binary size

How big are the CERN programs



Cacheline access frequency evaluated by sorting cachelines by their accesses

Thus a binary working set size measurement



Micro-kernels establish penalties and validate performance events

- To be useful for performance analysis, events must be able to be assigned a penalty
 - Otherwise it cannot be used to decompose cycle usage
- Do not assume you know what an event counts. You are almost certainly wrong
 - Documentation is always more vague than it appears
- Consider: 0xA0 L3_CACHE_RD L3 cache read
 - What does that include?
 - Ld, rfo, hw prefetch, sw prefetch, speculative + retired, secondary misses
 - If anything but retiring demand loads is included.
 - It is basically useless as it cannot be assigned a penalty

Micro-kernels establish penalties and validate performance events

- Data access latencies are established with linked list walker
 - `While(count < limit){count++;p=*p;}`
 - Need to worry about how to deal with hw prefetchers and paging
- Triads make good BW tests
 - w/wo RFO or streaming store
- Branch prediction, BTB capacity and ifetch are a bit trickier
- <https://github.com/David-Levinthal/gooda/tree/master/gooda-analyzer/kernels>
- Kernels produce precisely predicted counts for perf events
 - Validate that they count what they should
 - Harder to prove they do not count what they shouldn't
 - Run all events of interest against all kernels to look for errors

Micro-kernels establish penalties and validate performance events

```
levinth@ubuntu18-2:~/gooda/gooda-analyzer/kernels$ ls
```

| | | |
|---------------------|-----------------------------------|--------------------|
| arm | conditional_jump | pl_transition |
| bb_exec | dram_core_scan.sh | power |
| blocked_load | event_validation_instructions.txt | prefetch |
| br_multi_core | fourby.sh | README |
| c2c_core_scan.sh | gather | snoop_test |
| c2c_core_scan_ht.sh | inst_bw | sources_att |
| call_tests | latency_noarch | uncond_jump_rdm |
| callchain | mem_bw | unconditional_jump |
| cond_br_str | mem_latency | validate |

```
levinth@ubuntu18-2:~/gooda/gooda-analyzer/kernels$ ls arm
```

```
cond_br_str matrix_mult mem_bw mem_latency uncond_jump_rdm
```

Gooda has lots of documentation

```
levinth@ubuntu18-2:~/gooda/gooda-analyzer/docs$ ls
cycle_accounting_and_gooda.pdf          Micro-architecture.pdf
CycleAccountingandPerformanceAnalysis.pdf  old_stuff
Driving_the_Gooda_visualizer.pdf        useful_links.txt
GoodaInternals.pdf
levinth@ubuntu18-2:~/gooda/gooda-analyzer/docs$ ls old_stuff/
17775.pdf          kernel_cycles.pdf
17776.pdf          Loops.pdf
18027.pdf          NHM_NUMA.pdf
29672_itanium.pdf  Perf_analysis_Nehalem_intro.pdf
HPC_Perf_analysis_Xeon_5500_5600_intro.pdf
performance_analysis_guide.pdf
```