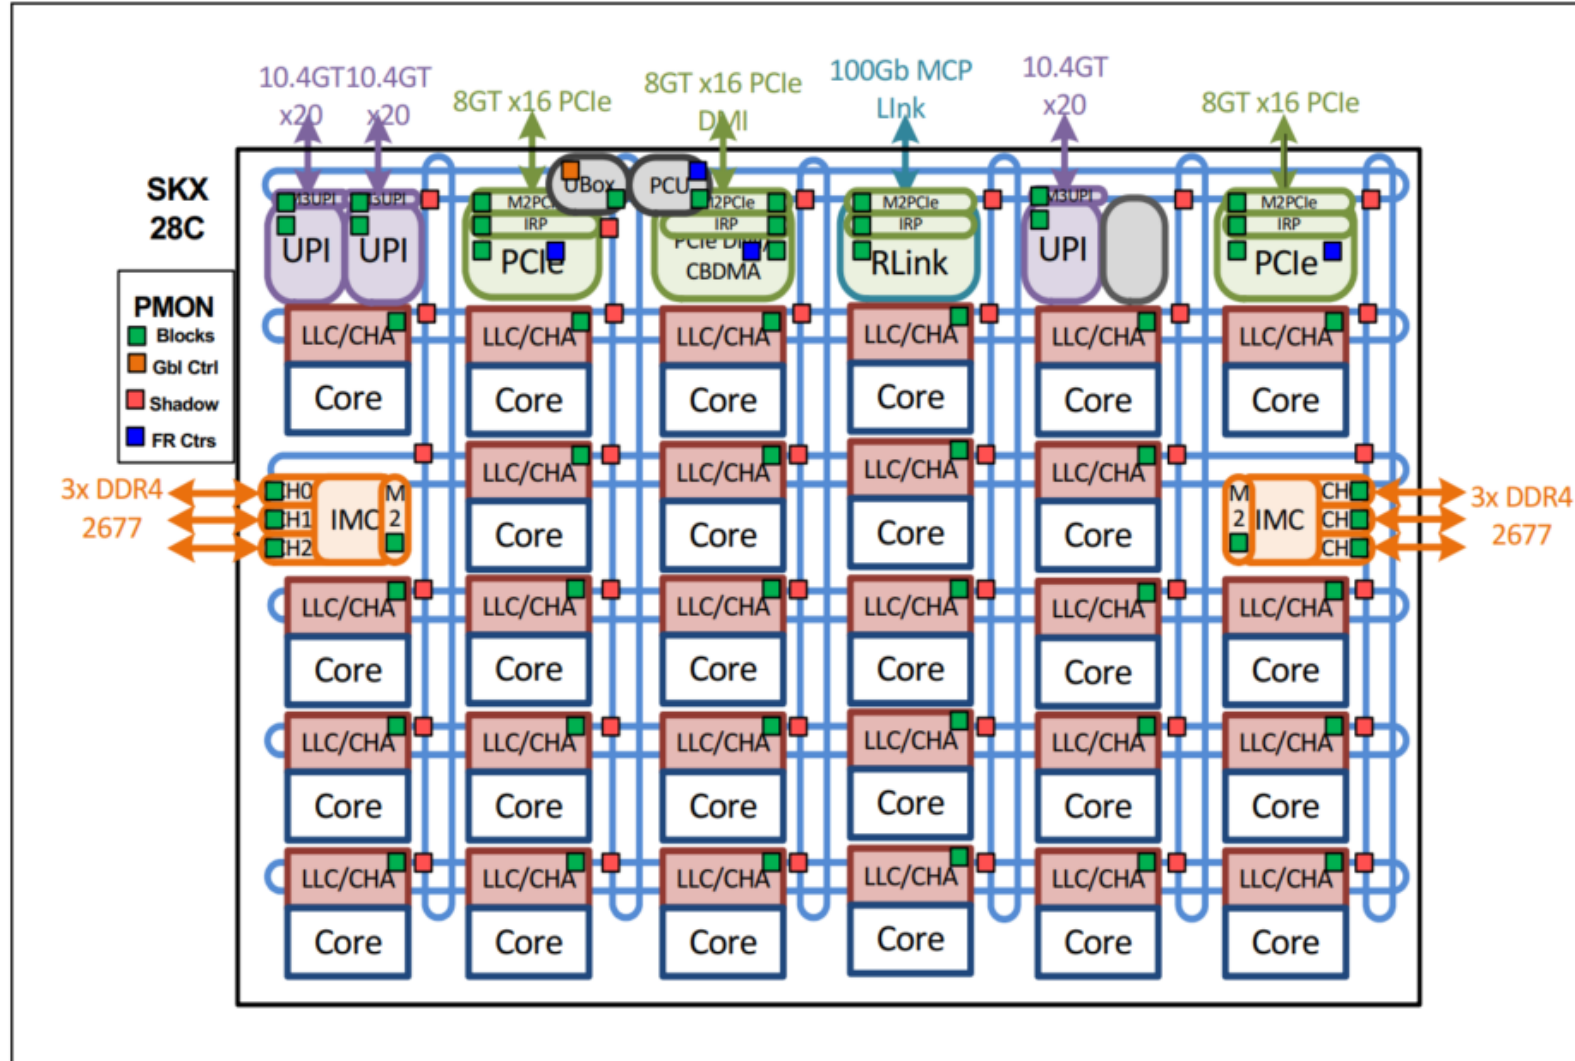


# SKX uncore monitoring

David Levinthal

See uncore PMU pdf 336274

**Figure 1-1. Intel® Xeon® Processor Scalable Memory Family - Block diagram for a 28C part**



# Basic Units

- Mesh made up of 2-dimensional structure of rings
  - 4 rings: AD(address) BL(Block/data) AK(Acknowledge) IV(Interrupts)
- CMS common mesh stop (aka common ring stop)
  - Home agent and Cbox combined in SKX and called CHA
  - Cache box and core interact with CMS
  - Cache box: unit of L3 cache/mesh stop. Lines associated by hash encoding of physical address, so there is no relation between the Cbox and the core at a single mesh stop, other than the common mesh stop (CMS)
  - See sec 2.2.7.1
- IMC: Integrated Memory Controller
  - M2M mesh to memory controller traffic management
- UPI
  - M3UPI mesh to UPI link Layer traffic management
- PCIe
  - IIO/M2PCIe is mesh to PCIe traffic manager
  - IRP coherency in IIO traffic
- Ubox (system control)
- PCU (Power control Unit)

# Uncore PMU resources

The general performance monitoring capabilities of each box are outlined in the following table.

**Table 1-8. Per-Box Performance Monitoring Capabilities**

Box	# Boxes	# Counters/ Box	# Queue Enabled	Packet Match/ Mask Filters?	Bit Width
CHA	up to 28	4	1	Y	48
IIO	up to 6 between C, P and M flavors	4 (+1) per stack (+4 per port)	0	N	48
IRP	up to 6	2	4	N	48
IMC	up to 2 (each with up to 3 channels)	4 (+1) (per channel)	4	N	48
Intel® UPI	up to 2 (2 or 3 links)	4 (per link)	4	Y	48
M3UPI	up to 2 (2 or 3 links)	3 (per link)	1	N	48
M2M	up to 2	4	1	Y	48
PCU	1	4 (+2)	4	N	48
UBox	1	2 (+1)	0	N	48

The programming interface of the counter registers and control registers fall into two address spaces:

- Accessed by MSR are PMON registers within the CHA units, IIO, IRP, PCU, and U-Box, see [Table 1-10](#).
- Access by PCI device configuration space are PMON registers within the IMC, Intel UPI and M3UPI units, see [Table 1-11](#).

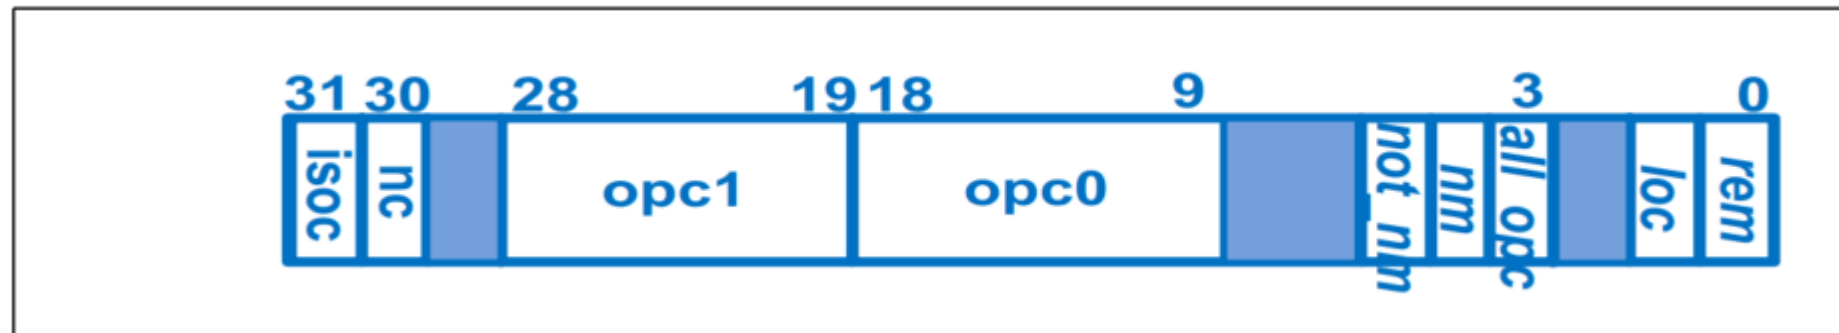
# CMS terminology

- Traffic (2.2.7.1)
  - Rx receive
  - Tx transmit
  - RxC receive from CMS
  - RxR receive from ring
  - VN0/1 virtual channel 0/1
- IA: core
- Queues (2.2.7.2)
  - IRQ: request from core (.IA events that were .IRQ events)
  - IPQ: ingress probe/remote socket snoop from UPI LL
  - ISMQ: response queue
  - PRQ: IIO request (.IO events that were .PRQ events)
  - RRQ: remote read queue
  - WBQ: writeback queue
- TOR: Table of Requests, tracking CHA transactions

# CHA event programming

- Basic event programming similar to core event programming
  - Event code, umask, threshold, invert, edge, etc
  - Emon threshold syntax is “:t=value”
- Two additional filters can be applied
  - Emon syntax is “:filter1=hexcode” or “:filter0=hexcode”
  - filter0 selects LLC mesi states (table 2-54) for LLC lookup events
  - filter1 selects request opcode (OR of up to 2 opc), local/remote etc
    - TOR\_occupancy/TOR\_inserts
    - Bottom byte = 0x33 is local or remote/opcode filter enable/cacheable or non cacheable

Figure 2-4. CHA PMON Filter 1 Register



# CHA events

## home agent, Core and CMS combined

- 1199 total events Some duplication to support old style home agent names
- TOR inserts and occupancies by request type / opcode matching
- Requests by MESI, Local, Remote, reads, writes, snoops, cross core snoops
- Directory lookups
- CMS interaction by receive, transmit
  - IPQ, RRQ, WBQ, PRQ
- LLC victimization
- IMC writes (memory BW per core?)
- Vert/Horiz traffic breakdown

# TOR Occupancy LLC miss event examples

- Note: these are monitoring the cache box requests from all cores not the core at the CMS
- UNC\_CHA\_TOR\_OCCUPANCY.IA\_MISS:filter1=0x40433
  - 0x33 is is local or remote/opcode filter enable/cacheable or non cacheable
  - 0x40400 is actually 202 in bit 18:9 (opc0)
    - Corresponds to demand data rd (see table 3.1)
  - OPC1 is not used
  - Counts TOR occupancy of demand data rd that miss the LLC
- UNC\_CHA\_TOR\_OCCUPANCY.IA\_MISS:filter1=0x12D40433
  - 0x33 is is local or remote/opcode filter enable/cacheable or non cacheable
  - 0x40400 is demand data rd in opc0
  - 0x12D00000 is 0x25A in bits 28:19 ie opc1
    - LLC data prefetch (from 3-1 sheet 2)
  - Counts occupancy of demand data rd OR prefetch data rd
    - ie all data requests from the core that miss the LLC



# Tor Occupancy LLC miss examples cont.

- UNC\_CHA\_TOR\_OCCUPANCY.IA\_MISS:filter1=0x12CC0233
  - 0x33 is is local or remote/opcode filter enable/cacheable or non cacheable
  - 0x12CC0200 is
    - 0x202 in opc0 or demand code rd from table 3-1
    - 0x259 in opc1 or LLcPrefCode from table 3-1 sheet 2
  - Counts demand code rd OR LLC prefetch code
    - ie all code requests from the core that miss the LLC
- UNC\_CHA\_TOR\_OCCUPANCY.IA\_MISS:filter1=0x12C40033
  - 0x12C40000 is
    - 0x200 in opc0 demand data RFO
    - 0x258 in opc1 LLcPrefRFO
  - Counts demand RFO OR LLC prefetch RFO
    - ie all reads for ownership exclusive state that miss the LLC

# Tor Occupancy LLC miss examples cont.

- UNC\_CHA\_TOR\_OCCUPANCY.IA\_MISS:filter1=0x12C40031
  - Same as before (ie all RFO)
  - Except 0x31 is remote target only, cacheable, non cacheable, opc enabled
  - Ie LLC RFO miss remote source

# IMC (416 events)

- Integrated memory controller can monitor traffic to and from dram
- Most basic events are total line reads and writes
  - UNC\_M\_CAS\_COUNT.RD
  - UNC\_M\_CAS\_COUNT.WR
- And the read and write queue occupancy and queue inserts
  - UNC\_M\_RPQ\_INSERTS
  - UNC\_M\_RPQ\_OCCUPANCY
  - UNC\_M\_WPQ\_INSERTS
  - UNC\_M\_WPQ\_OCCUPANCY
  - Note: occupancy is evaluated at IMC clockticks

# M2M IMC to mesh (457 events)

- Monitoring multi socket directory
  - UNC\_M2M\_DIRECTORY\_LOOKUP.\*
  - UNC\_M2M\_DIRECTORY\_UPDATE.\*
  - UNC\_M2M\_DIRECTORY\_HIT.\*
  - UNC\_M2M\_DIRECTORY\_MISS.\*
- Tracker activity
- WPQ activity
- Vertical and horizontal ring activity by ring
- 419 experimental events, 38 standard events

# UPI link layer (158 events)

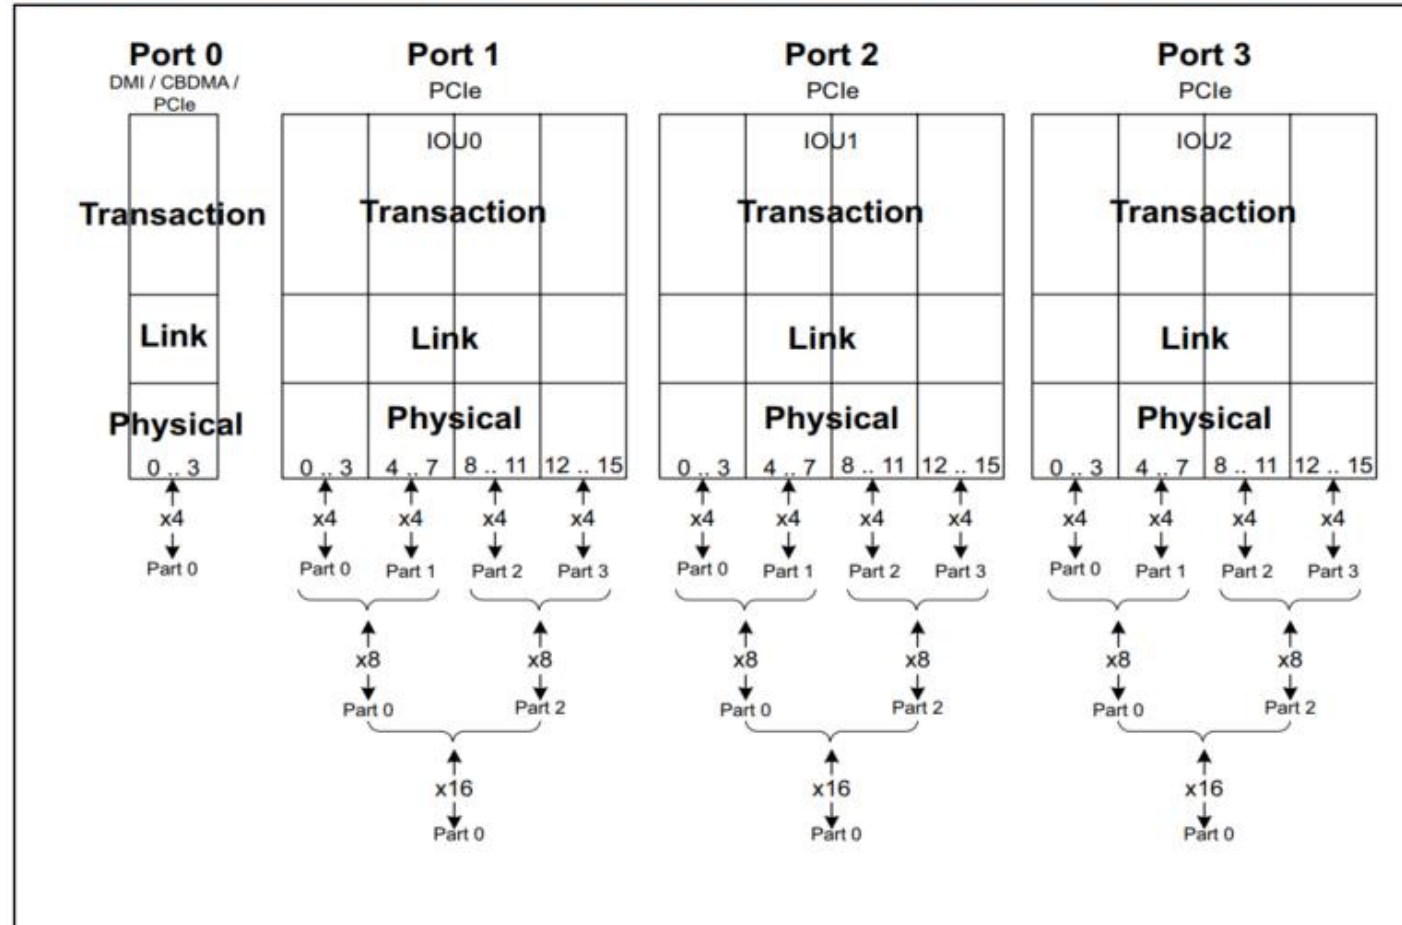
- Intersocket link layer traffic monitored (QPILL/UPILL)
  - Interaction with mesh monitored by m3upi
- UPI BW
  - `UNC_UPI_RxL_FLITS.ALL_DATA` (received flits)
  - `UNC_UPI_TxL_FLITS.DATA` (transmitted flits)
  - 9 flits needed to move a \$line so multiply by 64/9 to get bytes
- Link in various power mode
  - examples shows transmit, similar expression for receive
  - `UNC_UPI_RxLO_POWER_CYCLES/ UNC_UPI_CLOCKTICKS` (full power )
  - `UNC_UPI_TxLOP_POWER_CYCLES/ UNC_UPI_CLOCKTICKS` (low power)
  - `UNC_UPI_L1_POWER_CYCLES/ UNC_UPI_CLOCKTICKS` (shutdown Rx and Tx)
- Inserts/Occupancy per UPI slot (latencies)
- UPI headers opcode matching
- Cycles with no credits in flowQ (not quite sure what these are)
- Blocked receive queue cycles by cause

# M3UPI Mesh-UPI interface (629 events)

- Transmit and receive flowQ occupancies and inserts for VN0/1
  - Request, snoop, response and WB for AD and BL Rings
- Cycles FlowQ not empty
- Arbitration events for Tx, for AD and BL
  - Failed, Speculative arb for new, no credit, credit available
- Arbitration events for Rx, for AD and BL
  - “lost” and “can’t” arbitrate
- Breakdown by Vert and Horiz by ring of transmit (and some ring to ring)
  - Occupancy, inserts, cycles\_full

# PCIe and IIO monitoring structure

Figure 2-8. How IIO BW/TXNs are rolled up relative to width of attached Device(s)



Events such as DATA\_REQ\_OF\_CPU.MEM\_READ and TXN\_REQ\_BY\_CPU.PEER\_WRITE are further subdivided into '.PART's. For instance, if this is a IIO stack supporting up to x16 PCIe and a x16 Card is plugged into the stack, .PART0 will contain all the relevant information. If 2 x8 Cards are plugged into the stack, .PART0 will contain the contribution of the first x8 Card's traffic and .PART2 will contain the contribution of the second x8 Card's traffic. etc.

# PCIe and IIO monitoring with free running MSRs

- The PCIe monitoring has a unique feature of monitoring the input and output bandwidth in free running counters grouped by 4 lanes

**Table 1-11. Free-Running IIO Bandwidth Counters in MSR space**

	Port 3 BW In	Port 2 BW In	Port 1 BW In	Port 0 BW In	Port 3 BW Out	Port 2 BW Out	Port 1 BW Out	Port 0 BW Out
CBDMA	0x0B03	0x0B02	0x0B01	0x0B00	0x0B07	0x0B06	0x0B05	0x0B04
PCIe 0	0x0B13	0x0B12	0x0B11	0x0B10	0x0B17	0x0B16	0x0B15	0x0B14
PCIe 1	0x0B23	0x0B22	0x0B21	0x0B20	0x0B27	0x0B26	0x0B25	0x0B24
PCIe 2	0x0B33	0x0B32	0x0B31	0x0B30	0x0B37	0x0B36	0x0B35	0x0B34
MCP 0	0x0B43	0x0B42	0x0B41	0x0B40	0x0B47	0x0B46	0x0B45	0x0B44
MCP 1	0x0B53	0x0B52	0x0B51	0x0B50	0x0B57	0x0B56	0x0B55	0x0B54

**Table 1-12. Free-running IIO Utilization Counters in MSR space**

	Port 3 Util In	Port 2 Util In	Port 1 Util In	Port 0 Util In	Port 3 Util Out	Port 2 Util Out	Port 1 Util Out	Port 0 Util Out
CBDMA	0x0B0E	0x0B0C	0x0B0A	0x0B08	0x0B0F	0x0B0D	0x0B0B	0x0B09
PCIe 0	0x0B1E	0x0B1C	0x0B1A	0x0B18	0x0B1F	0x0B1D	0x0B1B	0x0B19
PCIe 1	0x0B2E	0x0B2C	0x0B2A	0x0B28	0x0B2F	0x0B2D	0x0B2B	0x0B29
PCIe 2	0x0B3E	0x0B3C	0x0B3A	0x0B38	0x0B3F	0x0B3D	0x0B3B	0x0B39
MCP 0	0x0B4E	0x0B4C	0x0B4A	0x0B48	0x0B4F	0x0B4D	0x0B4B	0x0B49
MCP 1	0x0B5E	0x0B5C	0x0B5A	0x0B58	0x0B5F	0x0B5D	0x0B5B	0x0B59



# PCIe and IIO monitoring with PMU

- Dominant IIO events are defined as matrices
- Event codes 83,84,c1,c2
  - Data\_req\_of/by\_cpu, TXN\_req\_of/by\_cpu
    - “Of” or “by” selects send vs receive
    - Data or TXN selects count of dwords or transaction requests
  - Read vs write
  - 4 “targets”
    - Memory (MMIO or dram depending on “by\_cpu” vs “of\_cpu” event type)
    - Peer (peer to peer)
    - IO (pcie IO space)
    - CFG (PCIeCFG space)
  - 6 PCIe partitions
    - part0-3
    - VTD0-1
  - Data\_req\_of\_cpu was called UNC\_IIO\_PAYLOAD\_BYTES\_IN
    - IO to IO transfer

# Measuring memory bandwidth

- Memory bandwidth is a time average count and can be measured at the memory controller
  - Read bw =  $64 * \text{UNC\_M\_CAS\_COUNT.RD} / \text{UNC\_M\_CLOCKTICKS}$
  - Write bw =  $64 * \text{UNC\_M\_CAS\_COUNT.WR} / \text{UNC\_M\_CLOCKTICKS}$
  - total bw =  $64 * \text{UNC\_M\_CAS\_COUNT.ALL} / \text{UNC\_M\_CLOCKTICKS}$

# Occupancy and insert events

- The ratio of occupancy/inserts is an average latency or time in the queue (aka Little's Law)
- There are 37 occupancy events with multiple umasks in the event files
  - Some support setting multiple umask bits
    - Result is the occupancy/inserts for the 'OR' of the umasks
  - Some do not
    - 'OR of umasks does not produce occupancy/inserts of the OR
  - Occupancy events that only have hardcoded umasks can't be ORed. ('00001000', '00000010', etc)
  - Occupancy events that use the common notation where each subevent has it's own bit with don't cares on the other bits. ('xxxx1xxxx', 'xxxxxx1x', etc)